

ECE2.414 – Advanced Memory Circuits and Systems

Assignment 1: Simulating Memory – Devices to Systems

Abhinav Venkata Kota
IIIT Hyderabad

Monsoon 2026

Part A — ngspice: 6T SRAM Read Margin

Central question: Does the bitcell actually read correctly, at 0.7 V and 85 °C?

A 6T SRAM cell was simulated using a 45 nm HP PTM BSIM4 model (`45nm_HP.pm`) in ngspice. The wordline (WL) ramps from 0 V to $V_{DD}=1.1$ V at $t = 1$ ns via a PWL source. The bitlines are precharged to V_{DD} and the cell stores $Q = 0$, $Q_B = V_{DD}$.

Task 1: $\Delta V(\text{BL}, \text{BLB})$ at $t = 2$ ns

Method: Transient analysis (`tran 5p 4n uic`) with initial conditions $v(Q) = 0$, $v(Q_B) = V_{DD}$. The differential voltage $\Delta V = v(\text{BLB}) - v(\text{BL})$ was computed as a separate vector and measured at $t = 2$ ns.

Parameter	Value
ΔV at $t = 2$ ns	729.86 mV
$q_{\max}$ (read disturb)	205.52 mV at $t = 1.0575$ ns

Table 1: Task 1 measurement results at 27 °C, $V_{DD} = 1.1$ V.

Comparison with reference: The measured $\Delta V = 729.86$ mV is $\sim 10.6\times$ larger than the 69 mV reference value from Lecture 1. The large differential indicates the cell reads correctly with ample margin. The 69 mV reference likely corresponds to a worst-case or minimum-margin scenario with different sizing or operating conditions.

Task 2: Plot $v(Q)$ and Effect of Doubling MA1/MA2

Base case ($W=0.16 \mu\text{m}$): During the read, the Q node is disturbed from 0 V up to $q_{\max} = 205.52$ mV at $t = 1.0575$ ns. This is the classic **read disturb** — current flows from the precharged BL through MA1 and MN1, pulling Q above ground. The cross-coupled inverter holds the data because MN1 ($W=0.20 \mu\text{m}$) is stronger than MA1 ($W=0.16 \mu\text{m}$), keeping the voltage divider favorable.

After doubling MA1/MA2 to $W=0.24 \mu\text{m}$:

Metric	$W=0.16 \mu\text{m}$	$W=0.24 \mu\text{m}$
ΔV (mV)	729.86	852.34
$q_{\max}$ (mV)	205.52	271.22
Cell ratio (W_{MA}/W_{MN})	0.80	1.20

Table 2: Task 2: Effect of doubling access transistor width.

With $W_{MA}=0.24\text{ }\mu\text{m}$ and $W_{MN}=0.20\text{ }\mu\text{m}$, the cell ratio drops to 1.2 — dangerously close to 1. The read disturb increases by 32% (271 mV vs 206 mV), pushing Q closer to the inverter trip point. At higher temperatures or lower V_{DD} , this configuration would likely **fail** (read upset), as the assignment predicts.

Task 3: ΔV vs V_{DD} Sweep

V_{DD} was swept from 1.1 V down to 0.6 V in 50 mV steps using a `.control` loop in ngspice.

V_{DD} (V)	ΔV at 2 ns (mV)
1.10	729.86
1.05	715.27
1.00	698.66
0.95	679.88
0.90	658.87
0.85	635.58
0.80	610.10
0.75	582.63
0.70	554.08
0.65	527.18
0.60	506.73

Table 3: Task 3: ΔV vs V_{DD} sweep results.

ΔV decreases monotonically with V_{DD} but remains well above 25 mV across the entire range (minimum $\Delta V = 506.73\text{ mV}$ at $V_{DD}=0.6\text{ V}$). With the given model and sizing, the bitcell maintains sufficient read margin even at 0.6 V. The 25 mV sense-amp offset threshold is never crossed in this sweep, indicating robust cell design for this voltage range.

Task 4: Repeat at $T = 85^\circ\text{C}$

Metric	27°C	85°C
ΔV (mV)	729.86	582.12
$q_{\max}$ (mV)	205.52	220.76
ΔV change	—	−20.2%
$q_{\max}$ change	—	+7.4%

Table 4: Task 4: Temperature comparison (27°C vs 85°C).

Analysis:

- **ΔV decreased by ~148 mV (20%):** Carrier mobility degrades with temperature ($\mu \propto T^{-1.5}$), reducing the drive current of the access transistors and slowing the bit-line discharge.
- **$q_{\max}$ increased by ~15 mV (7%):** The read disturb worsens because the pull-down NMOS (MN1) weakens more than the access transistor, shifting the voltage divider in favor of Q rising.
- **Which sets the failure?** The **sense-amp offset** moved more relative to its margin. While ΔV dropped 20%, it remains 582 mV — far above 25 mV. The real concern at

elevated temperature is the **read disturb** ($q_{\max}$ rising). If the cell were sized closer to failure, temperature-induced mobility degradation would push $q_{\max}$ past the inverter trip point before the sense-amp offset becomes the bottleneck.

Part B — CACTI: 2 MB SRAM L2 Sizing

Central question: How fast, how big, and how leaky is a 2 MB SRAM L2 at 45 nm?

CACTI 7 was used to model a 2 MB SRAM L2 cache at the 45 nm technology node. The configuration: 64 B block size, 8-way set-associative, 4 banks, ITRS-HP cell/peripheral types, ED² optimization, cycle-time-weighted objective.

Task 1: SRAM Baseline

Metric	Value
Access time	2.902 ns
Cycle time	2.652 ns
Dynamic read energy	0.793 nJ
Leakage power (bank)	562.6 mW
Area (H × W)	2.213 × 5.184 mm
Best Ndwl / Ndbl / Nspd	4 / 2 / 1

Table 5: Task 1: 2 MB SRAM L2 baseline at 45 nm, ED² optimized.

Task 2: Capacity Sweep (256 kB – 16 MB)

Capacity	Access (ns)	Energy (nJ)	Leakage (mW)
256 kB	2.417	0.658	90.3
512 kB	2.485	0.679	158.1
1 MB	2.634	0.717	293.7
2 MB	2.902	0.793	562.6
4 MB	3.531	0.930	1102.3
8 MB	4.473	1.220	2183.8
16 MB	6.362	1.742	4244.5

Table 6: Task 2: CACTI capacity sweep results.

Access time increases approximately linearly with $\log_2(\text{capacity})$. Each doubling adds ~ 0.2 – 0.5 ns, consistent with Amrutur & Horowitz’s “one gate delay per doubling” observation, though the exact scaling depends on the H-tree depth and bank organization.

Task 3: Optimization Objective Comparison

Objective	Access (ns)	Energy (nJ)	Ndwl	Ndbl
Pure delay	2.892	0.898	4	2
Pure area	3.355	0.797	2	2
ED ² (baseline)	2.902	0.795	4	2

Table 7: Task 3: Different optimization objectives for 2 MB SRAM.

The optimal organisation changes because each objective trades different physical resources: delay minimisation favours wider H-trees (higher N_{dwl}) and more parallel subarrays; area minimisation favours fewer splits (lower N_{dwl}) with denser mats; ED^2 balances both.

Task 4: Bitline Delay Hand-Check

From Task 1, CACTI reports a bitline delay of **0.407 ns**.

Using the hand-calculated estimate from Lecture 4:

$$\Delta t_{BL} \approx 0.38 \cdot R \cdot C \cdot L^2$$

The discrepancy arises because CACTI additionally models: (1) non-linear capacitance coupling between adjacent bitlines, (2) sense-amplifier offset and regenerative delay, (3) wire resistance and inductance in the H-tree routing, and (4) process-variation-aware timing margins. These effects are absent from the simplified $0.38 \cdot RC \cdot L^2$ formula.

Part C — NVSim: 2 MB STT-MRAM vs SRAM

Central question: Same capacity, same periphery — different bitcell. What changes?

NVSim is CACTI with the bitcell model extracted. Array organisation, peripheral search, and output format are identical; results are only as good as the `.cell` file. All runs below: **45 nm, HP, ReadEDP, Normal access, H-tree, 350 K, 2 MB, 64 B line, 8-way, 512-bit word** — identical to Part B. NVSim r131 patched for GCC 16 (`std::data` clash: `data` → `memData` in `main.cpp`:291,314,327,428), built binary 424 KB.

SRAM.cell: 146 F², 6T voltage-mode. `assign_STTTRAM.cell` (assignment spec, adapted to NVSim format): 54 F², aspect 2.0, $R_{on}=3\text{ k}\Omega$, $R_{off}=6\text{ k}\Omega$ (TMR 2:1), current-mode read $40\text{ }\mu\text{A}/0.25\text{ V}/25\text{ mV}$ sense, $145\text{ }\mu\text{W}$, write Reset/Set $200\text{ }\mu\text{A} \times 10\text{ ns}$, CMOS 6 F, $V_{drop}=0.15\text{ V}$.

Task 1: Side-by-Side 2 MB Cache

Reproduce:

```
1 cd NVSim
2 ./nvsim partC_SRAM_45nm_2MB.cfg
3 ./nvsim partC_STTTRAM_45nm_2MB.cfg
```

Metric (cache)	SRAM	STT-MRAM	Ratio
Area (total = data+tag)	6.393 mm ² (5.944+0.449)	3.590 mm ² (3.142+0.448)	0.56×
Read (hit) latency	0.850 ns	1.591 ns	1.87×
Write latency	0.520 ns	10.608 ns	20.4×
Read (hit) energy	0.575 nJ	0.763 nJ	1.33×
Write energy	0.531 nJ	0.435 nJ	0.82×
Leakage (total = data+tag)	3196 mW (2981+215)	425 mW (369+56)	0.13×

Table 8: Task 1: 6 required metrics side-by-side, 45 nm HP ReadEDP 2 MB 8-way. Ratios = STT/SRAM. Data/tag split in parentheses.

Detail (data array)	SRAM	STT-MRAM
Bank×Mat / Subarray	32×8 / 2×2	1×8 / 2×2
	64×256	256×128*
Area (eff.)	5.944 mm ² (83.4%)	3.142 mm ² (58.4%)
	2.935×2.025 mm	2.881×1.090 mm
Read lat.	753 ps	1449 ps
H-tree / Mat	467 / 287 ps	231 / 1218 ps
sense / prech.	3.5 / 52 ps	805 / 429 ps
Write lat. (pulse)	520 ps	10.608 ns (10 ns)
Read energy	561 pJ	665 pJ
H-tree / Mat	502 / 7.4 pJ	298 / 91.7 pJ
Write energy	518 pJ	390 pJ
Leakage	2981 mW (11.65/mat)	369 mW (23.09/mat)
Tag	0.449 mm ² , 251 ps, 14 pJ, 215 mW	0.448 mm ² , 1107 ps, 98 pJ, 56 mW

Table 9: Data/tag breakdown. *Fewer banks optimal for MRAM; tag similar, data dominates.

SRAM leakage is HP, no power-gating (CACTI would be lower); ratio still valid.

Task 2: Two Dramatically Worse, Two Dramatically Better

Worse: (1) **Write latency 20.4×** — STT needs 10 ns spin-transfer pulse to switch the free layer (Reset/SetPulse); SRAM flips cross-coupled inverters in ~ 65 ps charge. Device term: critical switching current I_{c0} and pulse width set by STT dynamics, not RC. (2) **Read latency 1.87×** (miss 4.4×) — current-mode sensing: $\Delta I \approx 40 \mu\text{A} \cdot (1/3\text{k} - 1/6\text{k}) \cdot R \rightarrow \sim 20$ mV vs >200 mV SRAM swing. TMR 2:1 $\rightarrow$ small ΔV , senseamp 805 ps vs 3.5 ps + precharge 429 ps vs 52 ps.

Better: (1) **Leakage 7.5×** **lower** — MTJ non-volatile, no 6T subthreshold nor cross-coupled contention; SRAM needs 256 mats (large area) $\rightarrow$ 2981 mW vs 369 mW despite higher per-mat leakage. (2) **Area 1.78×** **denser** — 1T-1MTJ 54 F² vs 6T 146 F² (2.7× cell); array $\sim 44\%$ smaller, shorter H-tree in one dimension. Read/write energy $\sim$ wash (H-tree dominates 300–500 pJ; cell read 7–15 pJ negligible at array level).

Task 3: TMR 2:1 $\rightarrow$ 3:1 ($R_{on}=4\text{k}$, $R_{off}=12\text{k}$)

Keep 40 μA read, 200 μA write, 54 F²:

```
1 # assign_STTRAM_TMRhigh.cell: 4000/12000 ohm
2 ./nvsim partC_STTRAM_TMRhigh_45nm_2MB.cfg
```

Metric	3k/6k (TMR 2:1)	4k/12k (TMR 3:1)	Delta
Total area	3.590 mm ²	3.580 mm ²	−0.28%
Read (hit) latency	1.591 ns	1.596 ns	+0.31% (+5 ps)
Write latency	10.608 ns	10.657 ns	+0.46% (+49 ps)
Read (hit) energy	0.763 nJ	0.761 nJ	−0.26%
Write energy	0.435 nJ	0.433 nJ	−0.46%
Leakage	425.2 mW	425.2 mW	0%

Table 10: Task 3: TMR sweep. Most moved: write latency +0.46%; all $<0.5\%$ — noise.

Most moved: read energy -1.9 pJ (noise). Device-paper TMR headline buys **margin/BER**, not speed/energy at array level: larger $\Delta R \rightarrow$ larger sense margin vs 25 mV offset, improves yield/ECC overhead and allows lower read current or smaller senseamp *if re-optimised* — but NVSim with fixed ReadPower/ReadCurrent and fixed senseamp keeps latency flat. Absolute R higher even slightly hurts (more IR drop). Array must co-optimize reference current/senseamp to monetise TMR.

Task 4: Halve Write Current $200 \rightarrow 100 \mu\text{A}$

```
1 # assign_STTRAM_halfCurrent.cell: Reset/Set 100uA
2 ./nvsim partC_STTRAM_halfCurrent_45nm_2MB.cfg
```

Metric	200 μA (6 F)	100 μA (6 F)	Delta
Total area	3.590 mm ²	3.590 mm ²	0%
Read (hit) latency	1.591 ns	1.571 ns	−1.26% (−20 ps)
Write latency	10.608 ns	10.588 ns	−0.19% (−20 ps)
Read energy	0.763 nJ	0.762 nJ	−0.13%
Write energy	0.435 nJ	0.433 nJ	− 0.46%
Leakage	425.2 mW	425.0 mW	−0.05%
Access width	6 F	6 F (3 F alt: same)	— (no change)

Table 11: Task 4: Halved write current. Most moved: write energy −0.46%; area 0% is artefact (fixed 54 F^2).

Mechanism — the single most important NVM coupling: $I_{\text{write}} \downarrow \rightarrow W_{\text{access}} \downarrow$ ($I_d \propto W$) $\rightarrow$ cell height/pitch $\downarrow \rightarrow$ CellArea (F^2) $\downarrow \rightarrow$ array area $\downarrow \rightarrow$ the density advantage you bought the technology for. Smaller TX also cuts bitline cap, leakage, and driver energy.

Why **zero area**: CellArea= 54 F^2 fixed dominates AccessCMOSWidth at 45 nm ($0.109 \mu\text{m}^2$ vs $6\text{ F} \approx 0.012 \mu\text{m}^2$). Model $\max(\text{CellArea}, \text{TX pitch}) \rightarrow 6\text{ F} \rightarrow 3\text{ F}$ still MTJ-limited. Si: $200 \mu\text{A}$ needs $\sim 8\text{--}10\text{ F} \rightarrow$ TX-limited; halving I would cut $\sim 30\text{--}40\%$ if CellArea derived. Artefact; ideal $\sim 0.07\text{--}0.08 \mu\text{m}^2$.